Count-level E0: the primary go/no-go gate fails

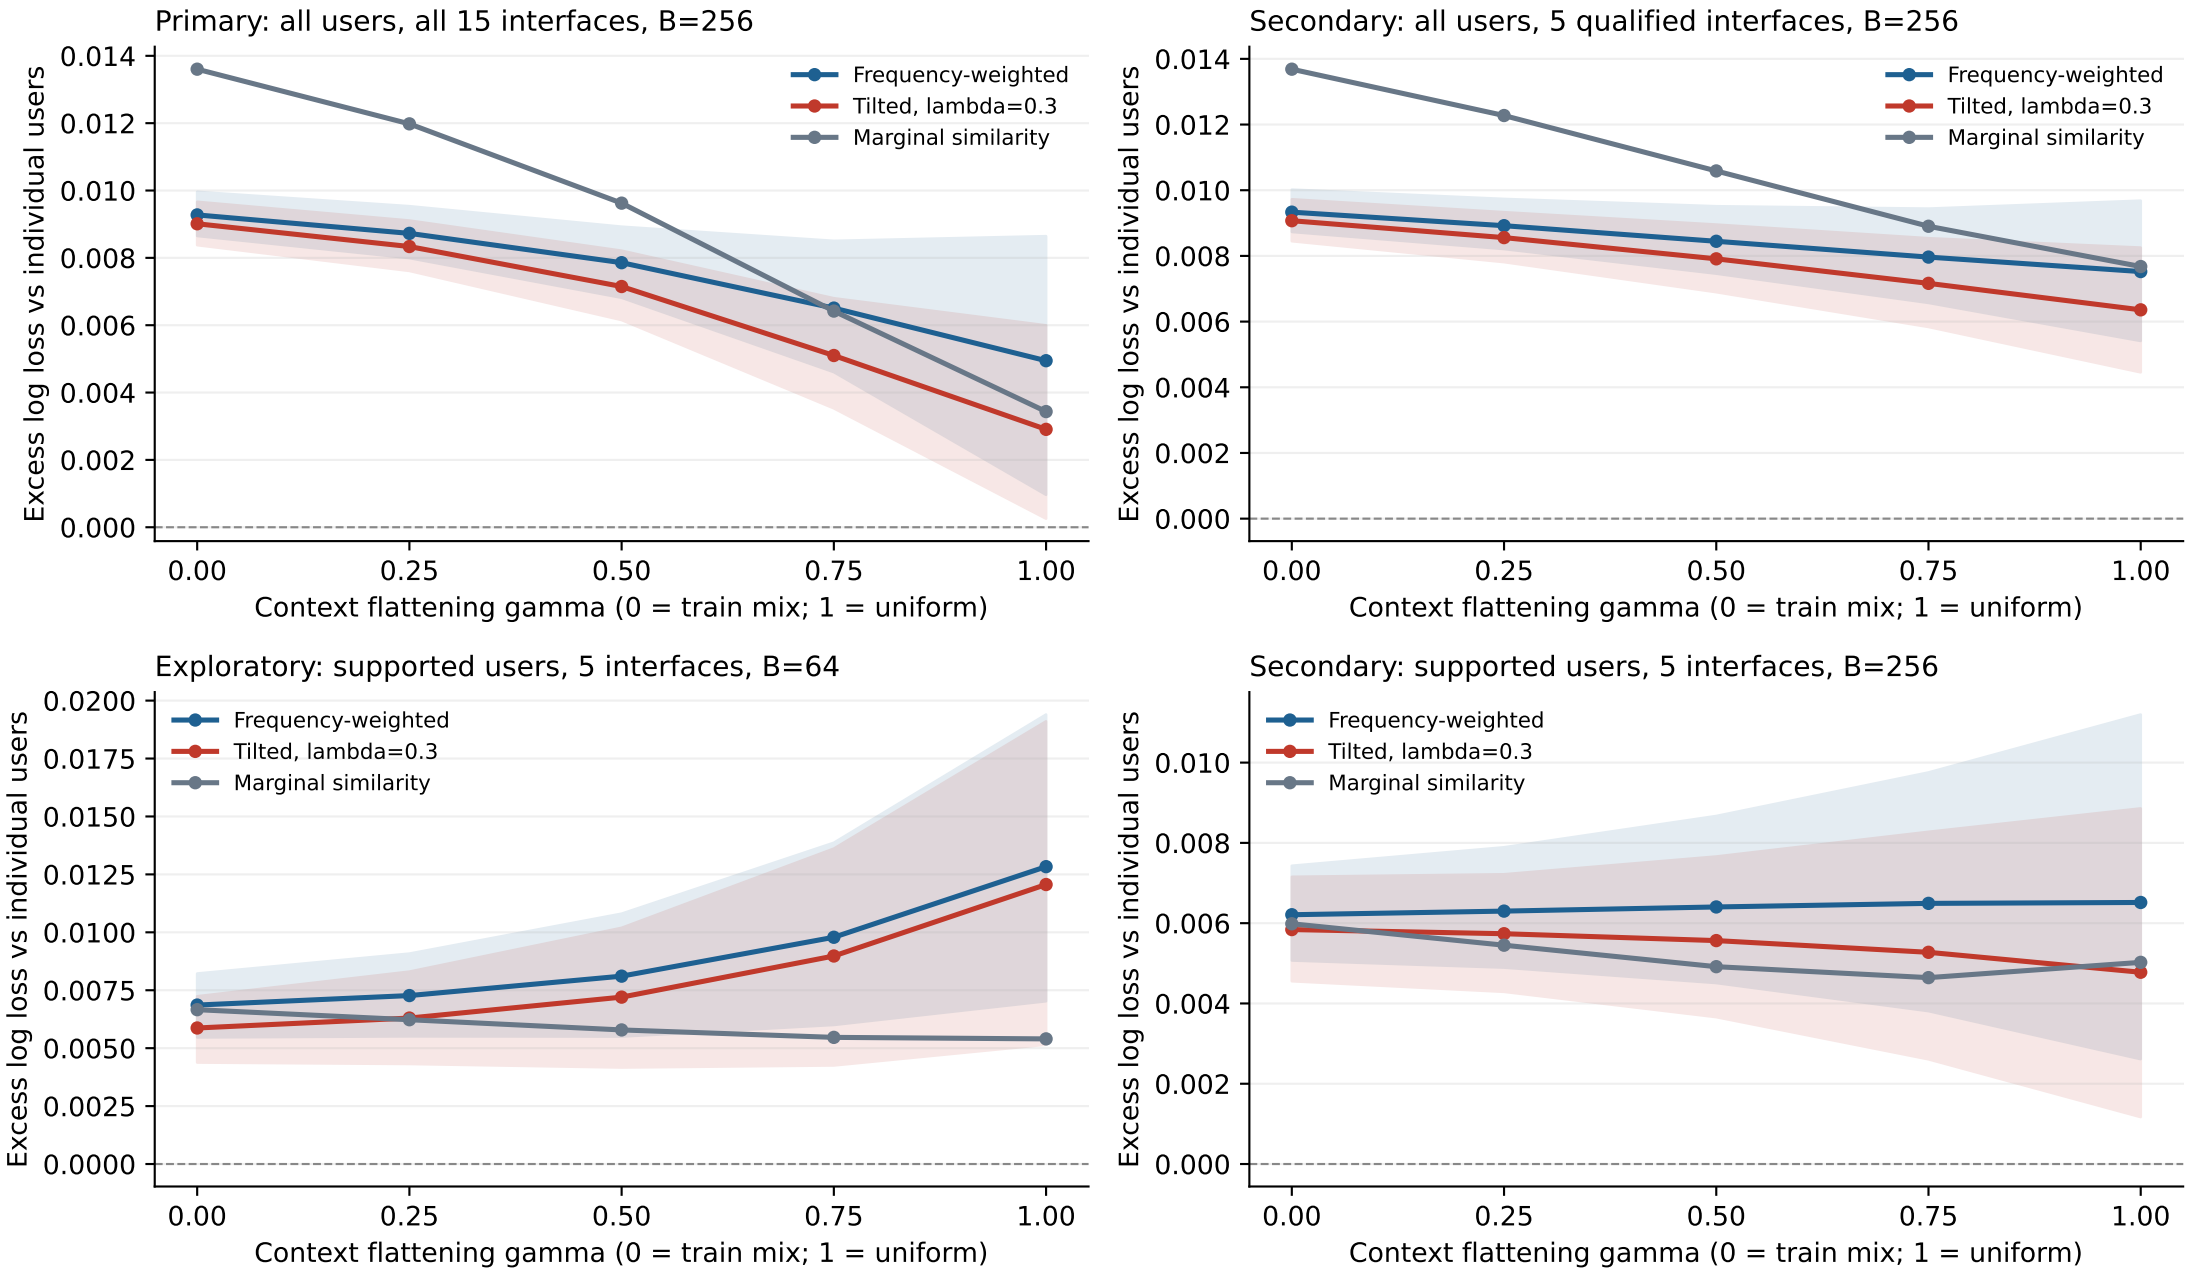

Lower excess loss is better. Shading: pointwise 95% user-bootstrap intervals. Fixed empirical context reweighting; exploratory test period.
Same bytes at each B; one behavioral initialization; capped heuristic. Qualified contexts chosen from training only: tabs 0, 1, 2, 4, 8.